

Ng Jun Xuan

FYP literature review unlinked (after word) turnitin

 FYP1 - Final Report

Document Details

Submission ID

trn:oid::3618:127860805

Submission Date

Feb 9, 2026, 9:47 PM GMT+8

Download Date

Feb 10, 2026, 12:30 PM GMT+8

File Name

FYP literature review unlinked (after word) turnitin.pdf

File Size

658.7 KB

40 Pages

8,033 Words

46,162 Characters

*% detected as AI

AI detection includes the possibility of false positives. Although some text in this submission is likely AI generated, scores below the 20% threshold are not surfaced because they have a higher likelihood of false positives.

Caution: Review required.

It is essential to understand the limitations of AI detection before making decisions about a student's work. We encourage you to learn more about Turnitin's AI detection capabilities before using the tool.

Disclaimer

Our AI writing assessment is designed to help educators identify text that might be prepared by a generative AI tool. Our AI writing assessment may not always be accurate (i.e., our AI models may produce either false positive results or false negative results), so it should not be used as the sole basis for adverse actions against a student. It takes further scrutiny and human judgment in conjunction with an organization's application of its specific academic policies to determine whether any academic misconduct has occurred.

Frequently Asked Questions

How should I interpret Turnitin's AI writing percentage and false positives?

The percentage shown in the AI writing report is the amount of qualifying text within the submission that Turnitin's AI writing detection model determines was either likely AI-generated text from a large-language model or likely AI-generated text that was likely revised using an AI paraphrase tool or word spinner.

False positives (incorrectly flagging human-written text as AI-generated) are a possibility in AI models.

AI detection scores under 20%, which we do not surface in new reports, have a higher likelihood of false positives. To reduce the likelihood of misinterpretation, no score or highlights are attributed and are indicated with an asterisk in the report (*%).

The AI writing percentage should not be the sole basis to determine whether misconduct has occurred. The reviewer/instructor should use the percentage as a means to start a formative conversation with their student and/or use it to examine the submitted assignment in accordance with their school's policies.

What does 'qualifying text' mean?

Our model only processes qualifying text in the form of long-form writing. Long-form writing means individual sentences contained in paragraphs that make up a longer piece of written work, such as an essay, a dissertation, or an article, etc. Qualifying text that has been determined to be likely AI-generated will be highlighted in cyan in the submission, and likely AI-generated and then likely AI-paraphrased will be highlighted purple.

Non-qualifying text, such as bullet points, annotated bibliographies, etc., will not be processed and can create disparity between the submission highlights and the percentage shown.





FACULTY OF
COMPUTING
& INFORMATICS

FYP01-DS-T2530-0860
Autonomous Driving Agent Training Using Deep
Reinforcement Learning in CARLA Simulator

242UC243LS

NG JUN XUAN

Abstract

This interim report systematically reviews the safety of Transformer-based Deep Reinforcement Learning (DRL) algorithms in the field of autonomous driving. Although reinforcement learning has shown potential in autonomous driving tasks, the safety performance differences among various reinforcement learning algorithms based on the Transformer architecture remain unclear. This study concludes that while the Soft Actor-Critic (SAC) algorithm offers higher overall safety in complex tasks, no single algorithm has an absolute advantage in all situations. Furthermore, this paper constructs a risk hierarchy through transitive reasoning and defines the turning intersection as the most complex scenario. Finally, the report proposes an implementation plan for Final Year Project 2 (FYP2) based on the results of this literature review. In Final Year Project 2, Transformer will be used as the main model for the SAC algorithm and applied to the intersection scenario.

Chapter 1: Introduction

1.1. Background of the Research

With the advent of the era of artificial intelligence and automation, many tasks no longer require human intervention (Frey & Osborne, 2017). Autonomous driving is one such trend, with most automakers integrating Level 4+ autonomous driving capabilities (Huynh & Nguyen, 2024, p. 1). Despite rapid development, autonomous driving has not yet been widely adopted. (Dosovitskiy et al., 2017, p. 7) point out that rule-based control methods are vulnerable in dynamic and complex environments, and safety is significantly reduced when the number of sensors is limited or when environmental changes are drastic. Most automakers not only need to consider the stability and cost of sensors, but the application of deep learning also relies on large amounts of labeled data, which requires significant manpower and time, resulting in high costs.

1.2. Problem Statement

Autonomous driving still faces many safety challenges in complex driving scenarios (Luo et al., 2025, p. 1). Although reinforcement learning has shown potential in autonomous driving tasks, the differences in safety performance among different RL algorithms using the Transformer architecture remain unclear. Therefore, it is difficult to determine which reinforcement learning algorithm is most suitable for developing and designing autonomous driving models.

1.3. Research Purpose

The purpose of this review is to comprehensively analyze and synthesize research on the safety of Transformer-based reinforcement learning algorithms in high-risk autonomous driving scenarios. Specifically, this review purpose is to:

- Classify commonly used Transformer architectures in autonomous driving tasks.

- Compare the safety performance of these algorithms (PPO, SAC, DDPG/TD3).
- Rank the risk levels of various scenarios based on data from all candidate papers.

This review not only identifies which algorithms offer the highest level of safety for autonomous driving but also serves as a foundational analysis, providing a reference for the design of FYP2.

1.4. Objectives of the Research

1. To systematically identify and categorize the most commonly used Transformer-based model and reinforcement learning algorithms in autonomous driving tasks.
2. To quantitatively compare the safety performance of reinforcement algorithms under comparable driving scenarios.
3. To rank the risk levels of different high-risk autonomous driving scenarios using a consistent set of safety metrics.

1.5. Research Questions

The main research question is under a Transformer-based policy network architecture, what are the differences in safety performance among different reinforcement learning algorithms in high-risk autonomous driving scenarios?

To answer this question, this study divides the main question into the following sub-questions:

- What are the differences in safety performance among different reinforcement learning algorithms in similar high-risk driving scenarios?
- What are some commonly used Transformer-based models in high-risk autonomous driving tasks?
- What are the differences between different high-risk scenarios?

1.6. Project Scope

The purpose of this study is to systematically evaluate the safety performance of different reinforcement learning algorithms under the Transformer architecture in high-risk autonomous driving scenarios.

1.6.1. Technical Scope

This literature review will analyze model architecture, algorithm performance, scenario risks, and safety metrics.

- This study uses Transformer-based models as the core component of the state encoder or policy network.
- This study primarily compares the safety performance differences of PPO, SAC, and DDPG/TD3 under the same Transformer architecture.
- This study will analyze several of the most common high-risk scenarios in the selected papers, including intersections, ramp merging, roundabouts, and highways.
- This study mainly uses quantitative safety metrics such as collision rate and success rate for evaluation.

1.6.2. Data Sources and Screening Scope

This literature review uses two sources: Google Scholar and IEEE Xplore. Google Scholar selects only the top 100 papers from 2019 to 2025, while IEEE Xplore selects all papers.

Only papers containing the Transformer architecture, specific reinforcement learning algorithms (PPO, SAC, DDPG/TD3), and safety-related metrics are considered. Papers involving non-autonomous driving scenarios, such as drones and driver state detection, will not be considered. Review articles and articles that are unavailable are not within the scope of this literature review.

1.6.3. Output and Contribution Scope

- Compare the safety performance of various algorithms under the same transformer model
- Rank the driving scenarios of each paper by risk level.
- Conclude the common types of transformers used in autonomous driving tasks.

1.7. Significance of the Research

This study, through a systematic literature review and comparative analysis of various metrics, explores the safety performance of different reinforcement learning algorithms under the Transformer architecture in high-risk autonomous driving scenarios. It has the following important theoretical and practical implications:

- This study constructs a comparative framework by collecting and extracting data from published papers, which helps to more objectively evaluate the differences in algorithm performance.
- This study also collects and extracts the performance of each algorithm in different scenarios to define and rank the risk levels of these scenarios.
- This study also reviews all types of Transformer-based models used and provides architectural options for future research.

These results will provide a reference for autonomous driving system developers in algorithm selection and optimization.

Chapter 2: Literature review

2.1. Introduction

Traditional rule-based methods require manually pre-setting driving strategies, which can lead to safety issues as scenarios become more complex (G. Zhao et al., 2025, p. 1). This chapter reviews the Transformer architecture and specifies reinforcement learning algorithms, explaining why the Transformer is suitable for application in the field of autonomous driving. Furthermore, this chapter provides definitions for various driving scenarios.

2.2. Systematic Literature Review Methodology

2.2.1. Search Strategies

- Database: Google Scholar, IEEE Xplore
- Years: 2019-2025
- Keyword: ("Transformer") AND ("Reinforcement Learning" OR "Deep Reinforcement Learning" OR DRL OR RL OR "Q-learning" OR "Policy Gradient") AND ("Autonomous Driving" OR "Autonomous Vehicle" OR "Self-driving" OR "Intelligent Vehicle" OR CAV) AND ("Collision Avoidance" OR "Time to Collision" OR TTC OR Safety OR "Near-miss" OR Accident OR Crash OR "Risk Assessment")

2.2.2. Include criteria

Included papers must meet all of the following criteria:

- The research must use a Transformer-based model as its core model or state encoder.
- Papers must cover at least one specified reinforcement learning algorithm, including PPO, SAC, or DDPG/TD3.
- Quantifiable safety assessment data must be provided.

2.2.3. Exclude criteria

Included papers cannot contain any of the following:

- All review articles, surveys, tutorials, or systematic reviews are excluded because they do not provide original experimental data.
- Research on non-autonomous driving, such as driver fatigue detection, is excluded. Research applied to off-road traffic scenarios (e.g., off-road vehicles) is also excluded.

2.2.4. Screening process




Figure 2.1 Screening Flow chart

2.3. Specific Applications of Transformer-RL in Autonomous Driving

In recent years, the Transformer architecture, with its self-attention mechanism, has demonstrated outstanding capabilities in handling long-range dependencies and complex relationships in modeling data, thus revolutionizing the field of autonomous driving (Doroudian & Taghavifar, 2024, p. 1). Transformer applications in reinforcement learning involve two embedding methods: one is using the Transformer as a state encoder, and the other is directly using the

Transformer as a policy/actor. These two application methods will be detailed below.

2.3.1. State Encoder Design (Multimodal Transformer)




Figure 2.2 Transformer as a state encoder

In driving tasks, multilayer perceptrons (MLPs) in reinforcement learning cannot directly process multimodal data; however, autonomous driving tasks rely on data from multiple sensors, such as camera and LiDAR point clouds.

The Transformer architecture overcomes these limitations through its self-attention mechanism (Vaswani et al., 2023). This mechanism computes relationships between sensor data in different formats, enabling the Transformer to correlate features from two modalities. For example, consider two sets of data for the Transformer: visual data and vehicle state. The agent sees a car rapidly approaching from the left front at an intersection (visual data), while simultaneously perceiving its own high speed and a slight leftward turn of the steering wheel (vehicle state). The attention mechanism allocates extremely high attention to "vehicle position" and "current speed," ignoring irrelevant roadside trees. The Transformer considers its current high-speed state (vehicle state) and the

rapidly approaching vehicle (visual features); if the vehicle does not correct its direction (steering state), a side collision will occur within seconds.

The DAD-RL framework proposed by (Chowdhury et al., 2024) employs a multimodal input design, with input data including historical vehicle trajectories and bird's-eye view (BEV) contextual information. The historical trajectory contains the position, heading, speed, and lane information of surrounding vehicles over the past 2.1 seconds, while the bird's-eye view includes a map of the drivable area and waypoints.

Besides the challenge of handling diverse data inputs, multilayer perceptrons (MLPs) also face the challenge of lacking the ability to process time-series data. Driving decisions in autonomous driving tasks inherently depend on both the current state and historical states. For example, if the agent only receives data from the current frame, such as a photo of a car rapidly approaching from the right at an intersection, but only one frame is provided, the agent cannot perceive the car's rapid approach and may even perceive it as stationary, potentially leading to catastrophic consequences.

While some models capable of processing time-series data, such as recurrent neural networks (RNNs) and long short-term memory networks (LSTMs), can serve as state encoders, their performance remains poor when processing excessively long time-series data. In contrast, Transformer utilizes a parallelizable self-attention mechanism, which enables it to efficiently process data with a large time span (Fan et al., 2024).

2.3.2. Policy Network Architecture (Transformer as Policy/Actor)




Figure 2.3 Transformer as a Policy/Actor

Compared to MLPs, Transformers are better suited as policy networks. MLPs lack explicit mechanisms for modeling temporal dependencies. Transformers, however, can analyze relationships between data spanning long time periods through self-attention mechanisms.

In their paper (Doroudian & Taghavifar, 2024), the authors integrated Transformers into the policy network of the PPO algorithm. The PPO Transformer (TPPO) achieved a success rate of 98% with a collision rate of only 2%, while the PPO without Transformers achieved a success rate of 42% and a collision rate of 12%. The safety of PPO with Transformers integrated is significantly improved compared to PPO without Transformers.

2.3.3. Difference between Transformer as Policy/Actor and Transformer as State Encoder

In methods using Transformer as a multimodal feature encoder, we can think of Transformer as a translator. Transformer "shreds" data of different formats and recombines it into a neat vector. This vector highlights what Transformer deems noteworthy and ignores unimportant

information. This vector is then fed to the reinforcement learning algorithm's MLP for decision-making.

Here are the benefits of using Transformer as a multimodal feature encoder:

- Transformer is computationally intensive, requiring only one computation, while the MLP can then use this "translated" data in a buffer for training.
- Developers can independently debug the "perception model" and the "decision model." This modular design reduces development complexity. For example, you can first train a Transformer to recognize objects, then train SAC to learn to drive.

(Chowdhury et al., 2024) uses a Spatio-Temporal Attention Encoder as the encoder and transforms the input data into a feature vector that the reinforcement learning algorithm can understand.

Transformer directly as a Policy/Actor Network: In this approach, Transformer is considered an "all-around brain." This method directly replaces the MLP in reinforcement learning algorithms with a Transformer model.

The following are the benefits of using the Transformer directly as the Policy/Actor Network:

- By using the Transformer directly as an Actor, it can receive historical data. This architecture without a "translator" avoids information that the actor considers important but the "translator" considers unimportant.

(Doroudian & Taghavifar, 2024) does not use the Transformer as a state encoder. Instead, it directly uses the Transformer as an actor of the policy/actor network (PPO).

2.4. Comparative Applications of RL Algorithms in Autonomous Driving Decision-Making

2.4.1. Application Research of PPO

In traditional reinforcement learning frameworks, policy gradient methods are extremely sensitive to the learning rate. On the one hand, an excessively small learning rate leads to slow parameter updates, significantly increasing the demand for training samples; on the other hand, an excessively large learning rate can have disastrous consequences. Since an agent's actions directly alter its subsequent state in reinforcement learning, large policy updates may cause the agent to enter previously unexplored unfavorable regions, resulting in a sharp decline in the quality of collected data and ultimately policy collapse. Imagine a novice driver discovering that turning right yields a high reward; in the next instance, they might sharply turn the steering wheel to gain more reward, leading to the collection of low-quality data.

To overcome this problem, (Schulman, Levine, et al., 2017) proposed Trust Domain Policy Optimization (TRPO). The core of TRPO lies in using Kullback-Leibler divergence to rigorously measure and constrain the difference between the old and new policies, thereby ensuring that each update is within a safe "trust domain." However, the brilliance of TRPO lies in its extremely high computational cost; the calculation of its second derivative (Hessian matrix) is extremely complex and difficult to implement. For example, to limit a student's sharp turns of the steering wheel, the instructor (KL divergence) calculates the similarity between the new and old actions and tells the student not to exceed a safe threshold.

Inspired by TRPO, Schulman, Wolski, et al. (2017) proposed Proximal Policy Optimization (PPO). The core innovation of PPO lies in the introduction of a "Clipped Surrogate Objective" which restricts the probability ratio of the new and old policies to a predefined interval (usually $[1 - \epsilon, 1 + \epsilon]$), thus achieving similar stability to TRPO with a

simpler first derivative approach. This mechanism not only significantly lowers the implementation threshold of the algorithm but also allows for multiple training iterations on the same batch of data, greatly improving sample efficiency. This is equivalent to the instructor adding a physical clip (Clipped Surrogate Objective) to the steering wheel, which restricts any sharp turns of the steering wheel.

Experiments show that PPO exhibits superior performance in both continuous control and discrete action space tasks (Schulman, Wolski et al., 2017). Compared to methods such as A2C, PPO maintains extremely high implementation flexibility while ensuring convergence speed and final performance. With its high stability and wide applicability, PPO has become one of the most highly regarded core algorithms in the fields of autonomous driving, robot control, and game AI.

2.4.2. Application Research of SAC

Soft Actor-Critic (SAC) is an offline deep reinforcement learning algorithm based on maximum entropy reinforcement learning and the stochastic actor-critic algorithm. It is particularly suitable for tasks with continuous action spaces, offering higher sample utilization and more stable training results (Haarnoja et al., 2018). SAC is a two-tube neural network (Haarnoja et al., 2018).

SAC uses one actor and two Q-function networks and can also be considered as critics. The actor's output is the parameters of the probability distribution of actions (usually the mean and standard deviation), for example, the throttle is $\mu = 0.5$, $\sigma = 0.28$. A small amount of noise is added to the standard deviation part to help the actor explore better actions. The critic part is similar to TD3, they both use two independent Q-function networks.

To put it simply, the two critics give the actor's behavior an evaluation. The two critics provide positive and negative feedback on the acceleration behavior. To prevent overly optimistic behavior, the actor will

listen to the more conservative opinion, which is the opinion of not accelerating.

2.4.3. Application Research of DDPG

Reinforcement learning methods for discrete action spaces, such as DQN, cannot be directly applied to continuous action spaces because the "maximization" operation of continuous actions requires complex optimization at each step, resulting in high computational cost and instability. DDPG combines the ideas of Deterministic Policy Gradient (DPG) and Deep Q-Networks (DQN) to propose an off-policy, model-agnostic actor-critic algorithm suitable for continuous action spaces (Lillicrap et al., 2019).

DDPG is based on the DPG framework, where:

- Actor: Outputs deterministic continuous actions.
- Critic: Evaluates the Q-values of state-action pairs.

By using a replay buffer and target networks, DDPG borrows the stable training mechanism of DQN, achieving efficient learning in continuous action spaces (Lillicrap et al., 2019).

2.4.4. Application Research of TD3

A problem exists in the DQN algorithm: Q-value overestimation. Since DDPG inherits from DQN, it also inherits this problem. This issue arises because the DQN critic overestimates the Q-value when evaluating actions due to noise, and actors tend to learn those overestimated actions, ultimately leading the policy to suboptimal solutions or even training failure.

To address this problem, TD3 was proposed. This algorithm introduces three main features:

- Delayed actor updates: This allows the critic to stabilize before the actor learns actions closer to the actual Q-value.
- Adding two independent critics: This ensures that the smallest Q-value is selected. While this might still lead the actor to learn actually underestimated actions, (Fujimoto et al., 2018) show that learning underestimated actions is safer than learning overestimated ones.
- Target policy smoothing: Adding a small amount of noise to the target action smooths the Q-value estimation and improves robustness.

2.5. Definition and Research of High-Risk Autonomous Driving Scenarios

2.5.1. Classification and Definition of High-Risk Scenarios

To cultivate robust decision-making in autonomous agents, exposure to high-risk scenarios during training is imperative. Analogous to the adage that calm seas do not produce skilled sailors, a benign training environment fails to equip agents with the necessary skills to handle unexpected real-world situations.

- Merging
 - Definition: A scenario where vehicles, without traffic signal guidance, need to observe, predict, and negotiate right-of-way with traffic flows from multiple directions.
 - Research Focus: Conflict resolution, intent prediction, and game theory decision-making.
- Roundabout
 - Definition: A multi-lane roundabout where vehicles need to merge, cycle, and exit within a continuous traffic flow.

- Research Focus: Cycle gap determination, lane keeping, and exit decision-making.
- Turn intersection
 - Definition: A turn intersection includes: unsignaled intersections, four-way intersections, left turns, and unprotected left turns. I categorize them all as turn intersections. It is a scenario where vehicles, without traffic signals, must observe, predict, and negotiate right-of-way with multi-directional traffic flows (including oncoming vehicles, pedestrians, and non-motorized vehicles).
 - Research Focus: Conflict resolution, intent prediction, and game theory decision-making.
- Highway
 - Definition: A high-speed, multi-lane environment. In some papers, the highway scenario may include overtake tasks.
 - Research Focus: Course deviation, maintaining distance, performance in severe weather.

2.5.2. Risk Assessment Index System

Each paper uses different metrics to measure the performance of its model and algorithm. Below is a summary of the most commonly used metrics across these papers. While some papers may use the same metrics, their measurement methods may vary slightly.

- Collision Rate: This is the most frequently used safety metric after data extraction. It calculates the proportion of collisions that occur in a test scenario.
- Collisions Time: This refers to the total number of collisions that occur when the autonomous vehicle completes a navigation task in each test scenario.

- **Success Rate:** This is the second most commonly used metric. This metric includes not only the collision-free task completion rate but may also include other factors, such as reaching the destination within the specified time.
- **Lane Departure Rate:** Commonly used to evaluate whether a vehicle can stably maintain its position near the center of the lane, avoiding excessive swaying or deviation. This metric is usually measured in centimeters (cm).

Chapter 3: Algorithm Performance Comparison

3.1. Introduction

This chapter aims to compare the safety performance of different deep reinforcement learning (DRL) algorithms in autonomous driving decision-making tasks through a systematic analysis of experimental data from existing literature. Although algorithms such as PPO, SAC, and DDPG/TD3 each have their theoretical focuses, their actual safety and effectiveness under the complex Transformer architecture still need to be verified through quantitative metrics. This chapter will collect and calculate the minimum, average, and maximum values of key metrics such as collision rate, collision time, and success rate for comparison. This chapter will focus on evaluating the safety performance of various mainstream algorithms based on average comparisons. Different methods are used for comparison of different metrics; a higher success rate is considered to indicate higher safety, and vice versa. Similarly, a lower collision rate and collision time are considered to indicate higher safety, and vice versa.

3.2. Comparative Analysis of PPO and SAC

Table 3.1 Performance Comparison between PPO and SAC

Paper title	PPO Performance	SAC Performance	Conclusion (Low performance < High performance)
Augmenting Reinforcement Learning With Transformer-Based Scene Representation Learning for Decision-Making of Autonomous Driving (Liu et al., 2024)	Collision rate: Min: 34% Mean: 51.2% Max: 64%	Collision rate: Min: 22% Mean: 35% Max: 52%	PPO < SAC
DAD-RL for Autonomous Decision-Making in Dynamic	Collision rate: Min: 34%	Collision rate: Min: 22%	PPO < SAC

Environment (Chowdhury et al., 2024)	Mean: 51.2%	Mean: 35%	
	Max: 64%	Max: 54%	

Table 3.1 (continued)

Efficient Maximum Entropy Reinforcement Learning Framework Guided by Expert Policies: An Application in Autonomous Driving (L. Gao et al., 2025)	Success rate:	Success rate:	PPO < SAC
	Min: 36%	Min: 65%	
	Mean: 37%	Mean: 68.5%	
	Max: 38%	Max: 72%	
Improving Reinforcement Learning with Expert Demonstrations and Vision Transformers for Autonomous Vehicle Control (Elallid et al., 2024)	Collision rate:	Collision rate:	PPO < SAC
	20%	6%	
	Success rate:	Success rate:	
	80%	94%	
SoftCTRL: Soft conservative KL-control of Transformer Reinforcement Learning for Autonomous Driving (Huynh & Nguyen, 2024)	Collision time:	Collision time:	PPO < SAC
	Mean: 20.87 ±1.14	Mean: 1.87 ±0.25	

Table 3.1 shows all the performance data that the papers include both PPO and SAC algorithms. In the data from these five papers, the SAC algorithm is significantly more secure than the PPO algorithm.

(Liu et al., 2024), PPO has a minimum collision rate of 34%, an average of 51.2%, and a maximum collision rate of 64%, while SAC has a minimum collision rate of 22%, an average of 35%, and a maximum collision rate of 52%.

This shows that SAC's collision rate is lower than PPO's collision rates. It is noteworthy in this table that the data from (Liu et al., 2024) and (Chowdhury et al., 2024) are very similar. This is because (Chowdhury et al., 2024) explicitly built upon (Liu et al., 2024)'s framework, using it as a baseline for comparison. This represents standard academic practice where new methods are evaluated against established benchmarks.

Regarding success rate, (L. Gao et al., 2025) shows that PPO's success rate is also lower than SAC's in terms of minimum, average, and maximum. (Elallid et al., 2024) used two metrics: in terms of collision rate, PPO (20%) was significantly higher than SAC (4%), while in terms of success rate, PPO (80%) and SAC (94%), both metrics showed that PPO was inferior to SAC. (Huynh & Nguyen, 2024) used a rarer metric: in terms of the number of collisions, PPO was nearly 11 times that of SAC.

3.3. Comparative Analysis of SAC and DDPG/TD3

Table 3.2 Performance Comparison between SAC and DDPG/TD3

Paper title	SAC Performance	DDPG/TD3 Performance	Conclusion (Low performance < High performance)
Improving policy training for autonomous driving through randomized ensembled double Q-learning with Transformer encoder feature evaluation (Fan et al., 2024)	Success rate: Min: 16.2% Max: 16.4%	Success rate: Min: 16.3% Max: 16.4%	SAC $\approx$ DDPG/TD3

Table 3.2 shows performance data for both the SAC and TD3 (DDPG) algorithms. In (Fan et al., 2024), the performance of SAC and TD3 is almost identical.

3.4. Comparative Analysis of Merging and Intersection Scenarios DDPG/TD3 and PPO

Table 3.3 Performance Comparison between DDPG/TD3 and PPO

Paper title	DDPG/TD3 Performance	PPO Performance	Conclusion (Low performance < High performance)
Autonomous Intersection Management via Prior-Enhanced Multi-Agent Constrained Decision Transformer (R. Zhao et al., 2025)	Collision rate: Min: 44.37% Mean: 58.98% Max: 69.86%	Collision rate: Min: 60.25% Mean: 66.95% Max: 77.99%	PPO < DDPG/TD3
Constraint-Guided Behavior Transformer for Centralized Coordination of Connected and Automated Vehicles at Intersections (Zhao et al., 2024)	Collision rate: Min: 66.73% Mean: 68.295% Max: 69.86%	Collision rate: Min: 64.265 Mean: 71.125% Max: 77.99%	PPO < DDPG/TD3

In (R. Zhao et al., 2025; Zhao et al., 2024), the average collision rates of DDPG/TD3 were 58.98% and 68.295%, respectively, while those of PPO were 66.95% and 71.125%. These results indicate that PPO performs worse than DDPG/TD3.

3.5. Comparative Analysis of Merging and Intersection Scenarios PPO, SAC and DDPG/TD3

Table 3.4 Performance Comparison between PPO, SAC and DDPG/TD3

Paper title	PPO Performance	SAC Performance	DDPG/TD3 Performance	Conclusion (Low performance < High performance)
Safety-constrained Reinforcement Learning with Interaction-aware for Decision-making of Autonomous Driving (Zhang et al., 2025)	Success rate: Min: 42% Mean: 51.33% Max: 58%	Success rate: Min: 33% Mean: 38.7% Max: 48%	Success rate: Min: 7% Mean: 18.7% Max: 37%	DDPG/TD3 < SAC SAC < PPO DDPG/TD3 < PPO

(Zhang et al., 2025) is the most direct in determining which algorithm is optimal because it directly compares the three. (Zhang et al., 2025) shows that PPO is the best, with an average success rate of 51.33%, followed by SAC with an average success rate of 38.7%, and finally DDPG/TD3 with an average success rate of 18.7%.

3.6. Analysis summary

Table 3.5 Comparative Performance Analysis of Reinforcement Learning Algorithms

No	Conclusion (Low performance < High performance)	Support paper
1	$PPO < SAC$	(Liu et al., 2024) (Chowdhury et al., 2024) (L. Gao et al., 2025) (Elallid et al., 2024) (Huynh & Nguyen, 2024)
2	$SAC \approx DDPG/TD3$	(Fan et al., 2024)

Table 3.5 (continued)

3	$PPO < DDPG/TD3$	(R. Zhao et al., 2025) (Zhao et al., 2024)
4	$DDPG/TD3 < SAC$	(Zhang et al., 2025)
5	$SAC < PPO$	(Zhang et al., 2025)
6	$DDPG/TD3 < PPO$	(Zhang et al., 2025)

Table 3.5 summarizes the performance comparisons of the PPO, SAC, and DDPG/TD3 algorithms in the main studies. This table reveals two contradictions. The first is the comparison between PPO and SAC. While most studies (5 supporting papers) agree that $PPO < SAC$, indicating that SAC has higher efficiency in autonomous driving tasks, a few studies (1 paper) reach the opposite conclusion: $SAC < PPO$. The second is the comparison between PPO and DDPG/TD3. Two papers support the claim that $PPO < DDPG/TD3$, while only one paper supports $PPO > DDPG/TD3$.

In conclusion, there are a total of 8 papers covering the PPO algorithm, 7 papers covering the SAC algorithm, and 3 papers covering the DDPG/TD3 algorithm.

This study, through a systematic literature review, finds that PPO, SAC, and DDPG/TD3 do not have an absolute superiority or inferiority in autonomous driving decision-making tasks. Their performance is highly dependent on task attributes, state representation complexity, reward function design, and environmental stochasticity. For example, PPO performs better in safety-constrained tasks requiring stable policy updates, while SAC is more advantageous in scene representation learning where exploration efficiency is crucial. Therefore, algorithm selection should be based on a thorough analysis of task characteristics.

Chapter 4: High-Risk Scenario Risk Assessment

4.1. Introduction

(Lu et al., 2025, p. 1) points out that autonomous driving tasks have multiple characteristics, including high vehicle density, diverse traffic participants, and complex traffic rules. These factors affect the safety and robustness of autonomous driving algorithms. This chapter summarizes and extracts data from all relevant papers and calculates their minimum, mean, and maximum values. Different comparison methods are used for different indicators. Higher collision rates, collision time and lane deviation are considered high risk, and vice versa. Higher success rates are considered low risk, and vice versa. At the last of this chapter, logical reasoning will be used to construct a comprehensive ranking of scenario risk levels.

4.2. Comparative Analysis of Merging and Intersection Scenarios

Table 4.1 Risk Assessment between Merging and Intersection

Paper title	Merging risk assessment	Intersection risk (Left turn) assessment	Conclusion (Low risk < High risk)
Efficient Maximum Entropy Reinforcement Learning Framework Guided by Expert Policies: An Application in Autonomous Driving (L. Gao et al., 2025)	Success rate:	Success rate:	Intersection (Left turn) < Merging
	Min: 38%	Min: 36%	
	Mean: 69.14%	Mean: 72%	
	Max: 92%	Max: 96%	

Table 4.1 compares the performance of autonomous driving in two high-risk scenarios—merging and intersections. As you can see, the average success rate at intersections is 72%, and at merging it's 69.14%, with little difference between the two. This indicates that in this study, the risks of these two scenarios are comparable, but merging is slightly more dangerous.

4.3. Comparative Analysis of Roundabout and Intersection

Table 4.2 Risk Assessment between Roundabout and Intersection

Paper title	Roundabout risk assessment	Intersection (Right turn) risk assessment	Conclusion (Low risk < High risk)
Improving Reinforcement Learning with Expert Demonstrations and Vision Transformers for Autonomous Vehicle Control (Elallid et al., 2024)	Collision rate: 6%	Collision rate: 32%	Roundabout < Intersection (Right turn)
	Roundabout risk assessment	Intersection risk assessment	Conclusion (Low risk < High risk)
Autonomous Driving Agents for Safe Road Junctions Through Deep Reinforcement Learning (Yun & Cao, 2024)	Collision rate: Min: 7.2% Mean: 14.12% Max: 20.47%	Collision rate: Min: 56.88% Mean: 62.28% Max: 65.51%	Roundabout < Intersection

Table 4.2 compares the performance of autonomous driving models in roundabout and intersection scenarios. (Elallid et al., 2024) shows that autonomous driving models are more prone to accidents in intersection scenarios, with a collision rate of 32% compared to 6% in roundabout scenarios. Beside, (Yun & Cao, 2024) also indicates that the intersection scenario is higher risk than roundabout.

4.4. Comparative Analysis of Intersection and Highway

Table 4.3 Risk Assessment between Intersection and Highway

Paper title	Intersection risk assessment	Highway risk assessment	Conclusion (Low risk < High risk)
Empirical Evaluation of Deep Learning based Autonomous Driving System with Real-Time Scene Classification Procedure (Prasath A & Navaz K, 2025)	Lane deviation error: Min: 6.47cm Mean: 9.77cm Max: 12.4cm	Lane deviation error: Min: 6.22cm Mean: 8.96cm Max: 13.45cm	Highway < Intersection
	Intersection risk (Left turn) assessment	Highway risk assessment	Conclusion (Low risk < High risk)
Pre-trained Transformer-Enabled Strategies with Human-Guided Fine-Tuning for End-to-end Navigation of Autonomous Vehicles (Hu et al., 2024)	Collision rate: Min: 0% Mean: 21.93% Max: 58.7%	Collision rate: Min: 0% Mean: 8.84% Max: 38.7%	Highway < Intersection (Left turn)

Table 4.3 contains two papers, each using a different metric: lane deviation error and collision rate. Both papers show that the risk is higher in intersection scenarios than in highway scenarios. In (Prasath A & Navaz K, 2025), the average lane deviation error for intersections is 9.77 cm, while the average lane deviation error for highways is 8.96 cm. (Hu et al., 2024) shows an average collision rate of 21.93% for intersections, compared to 8.84% for highway scenarios.

4.5. Comparative Analysis of Merging, Roundabout and Intersection

Table 4.4 Risk Assessment between Roundabout and Intersection

Paper title	Merging risk assessment	Roundabout risk assessment	Intersection (Left turn) risk assessment	Conclusion (Low risk < High risk)
Augmenting Reinforcement Learning With Transformer-Based Scene Representation Learning for Decision-Making of Autonomous Driving (Liu et al., 2024)	Collision rate: Min: 4% Mean: 39% Max: 100%	Collision rate: Min: 16% Mean: 40.56% Max: 100%	Collision rate: Min: 12% Mean: 32.66% Max: 54%	Merging < Roundabout Intersection (Left turn) < Roundabout Intersection (Left turn) < Merging
Deep Attention Driven Reinforcement Learning (DAD-RL) for Autonomous Decision-Making in Dynamic Environment (Chowdhury et al., 2024)	Collision rate: Min: 2% Mean: 31.25% Max: 100%	Collision rate: Min: 4% Mean: 34.17% Max: 100%	Collision rate: Min: 0% Mean: 25% Max: 54%	Merging < Roundabout Intersection (Left turn) < Roundabout Intersection (Left turn) < Merging

Table 4.4 contains two papers that use the same metric and include scenarios involving merging, roundabout, and intersection. (H. Liu et al., 2024)'s average collision rates are 39% for merging, 40.56% for roundabouts, and 32.66% for intersections. (Chowdhury et al., 2024)'s average collision rates are 31.25% for merging, 34.17% for roundabout, and 25% for intersection. Both papers show that roundabouts are the most dangerous, followed by merging, and lastly, intersections.

4.6. Comparative Analysis of Merging, Intersection and Highway

Table 4.5 Risk Assessment between Merging, Intersection and Highway

Paper title	Merging risk assessment	Intersection (Left turn) risk assessment	Highway risk assessment	Conclusion (Low risk < High risk)
Safety-constrained Reinforcement Learning with Interaction-aware for Decision-making of Autonomous Driving (Zhang et al., 2025)	Success rate: Min: 7% Mean: 44% Max: 94%	Success rate: Min: 12% Mean: 49.5% Max: 97%	Success rate: Min: 37% Mean: 60.75% Max: 100%	Intersection (Left turn) < Merging Highway < Intersection (Left turn) Highway < Merging

(Zhang et al., 2025) includes scenarios involving merging, intersections, and highways. The average success rate of the first paper was 44% for merging, 49.5% for intersection, and 60.75% for highway. This paper shows that merging is the most dangerous, followed by intersections, and lastly highways.

4.7. Analysis summary

Table 4.6 Summary of Scenario Risk Hierarchies from Primary Studies

No	Conclusion (low risk < high risk)	Support paper
1	<i>Intersection(Left turn) < Merging</i>	(L. Gao et al., 2025) (H. Liu et al., 2024) (Chowdhury et al., 2024) (Zhang et al., 2025)
2	<i>Roundabout < Intersection (Right turn)</i>	(Elallid et al., 2024)
3	<i>Roundabout < Intersection</i>	(Yun & Cao, 2024)
4	<i>Intersection (Left turn) < Roundabout</i>	(H. Liu et al., 2024) (Chowdhury et al., 2024)

Table 4.6 (continued)

5	<i>Highway < Intersection</i>	(Prasath A & Navaz K, 2025)
6	<i>Highway < Intersection (Left turn)</i>	(Hu et al., 2024) (Zhang et al., 2025)
7	<i>Merging < Roundabout</i>	(H. Liu et al., 2024) (Chowdhury et al., 2024)
8	<i>Highway < Merging</i>	(Zhang et al., 2025)

Table 4.6 summarizes the risk comparison findings extracted from all studies, which forms the basic logical elements for subsequently constructing an overall scenario hierarchy. Each entry follows a formal relational logic, with the left side representing perceived risk or complexity lower than the right side. As shown in the table, empirical evidence from recent literature.

Table 4.7 Logic Synthesis of Driving Scenario Complexity Hierarchies

Logic Components	Synthesized Hierarchy
$(I(l) < Rab) + (I(l) < Mrg) + (Mrg < Rab)$	$I(l) < Mrg < Rab$
$(H < I) + (H < I(l)) + (H < Mrg) + (I(l) < Mrg)$	$H < ((I(l))/I) < Mrg$
$(H < (I(l)/I)) < Mrg) + (I(l)) < Mrg < Rab)$	$H < (I(l)/I) < Mrg < Rab$
$(H < (I(l)/I) < Mrg < Rab) + (Rab < I(r)) + Rab < I$	$H < (I(l)/I) < Mrg < Rab < (I(r)/I)$

This chapter systematically constructs a hierarchical model of risk and complexity in autonomous driving scenarios through qualitative analysis of eight primary studies.

Based on the literature evidence listed in Table 4.6, this study extracts direct comparative data on relative risk levels under different traffic environments. The study finds that although different scholars have different research focuses, they show a high degree of consistency in core conclusions such as "highways

have the lowest complexity" and "intersection turning tasks have high complexity."

Through the derivation process of the logic components shown in Table 4.7, this study uses transitive inference to synthesize fragmented evidence into a complete hierarchical architecture. The final synthesized hierarchical model (Highway < Intersection (Left turn) / Intersection < Merging < Roundabout < Intersection (Right turn) / Intersection) clearly characterizes the gradient of task challenges faced by autonomous driving systems in the operating domain.

The scenario hierarchy model established in this chapter not only verifies the consistency of existing literature, but also provides theoretical support and scenario selection basis for the research on decision-making algorithms for high-complexity scenarios (such as right-turn intersections) in subsequent chapters.

Chapter 5: Transformer Architectures in Autonomous Driving

5.1. Introduction

In autonomous driving, there are various variants based on the Transformer, each designed to address specific shortcomings, such as training efficiency or suitability for different perception modalities. The following are some of the more common Transformer variants used in autonomous driving.

5.2. Vision Transformer

Table 5.1: Papers that use the Vision Transformer model

Paper title
A Preference-Based Multi-Agent Federated Reinforcement Learning Algorithm Framework for Trustworthy Interactive Urban Autonomous Driving (Lu et al., 2025)
Deep Reinforcement Learning for a Self-Driving Vehicle Operating Solely on Visual Information (Lu et al., 2025)
Efficient Maximum Entropy Reinforcement Learning Framework Guided by Expert Policies: An Application in Autonomous Driving (L. Gao et al., 2025)
Improving Reinforcement Learning with Expert Demonstrations and Vision Transformers for Autonomous Vehicle Control (Elallid et al., 2024)
Integrating Driving-Aware World Model With MPC for Autonomous Driving at Unsignalized T-Intersections (X. Zhang et al., 2025)
Vision Transformer for Learning Driving Policies in Complex and Dynamic Environments (Kargar & Kyrki, 2022)

Vision Transformer (ViT) is a model that directly applies the Transformer architecture to image classification tasks. It segments images into fixed-size patches, linearly embeds these patches into a one-dimensional sequence, and then feeds it into a standard Transformer encoder for processing. ViT performs well in multiple computer vision tasks and is widely used in the field of autonomous driving for end-to-end perception and decision-making tasks.

(Audinys et al., 2025) installed a camera in front of a virtual vehicle to capture 224×224 pixels RGB images. In order to enable the model to infer the vehicle's current motion state (such as speed and acceleration), the input was not a single image, but rather four consecutive frames stacked together as a set of inputs.

(G. Li et al., 2025) used the Swin Transformer, a model with a hierarchical structure and a shifted window mechanism compared to ViT. This design allows the model to recognize both macroscopic objects and subtle local details in images, similar to a CNN, making it well-suited for object detection and semantic segmentation tasks. The image is divided into multiple non-overlapping local windows (e.g., 4×4 or 7×7 patches), and self-attention computation is performed only within each window.

5.3. Decision Transformer

Table 5.2: Papers that use the Decision Transformer model

Paper title
Autonomous Intersection Management via Prior-Enhanced Multi-Agent Constrained Decision Transformer (R. Zhao et al., 2025)
Deep reinforcement learning navigation via decision transformer in autonomous driving (Ge et al., 2024)

Decision Transformer (DT) is a model that uses a transformer architecture to decision-making in reinforcement learning (RL). Contrary to policy gradient-based reinforcement learning, DT views the sequential decision-making problem as an autoregressive sequence modeling problem. It is a direct predictor of the future action sequence based on a sequence of states, actions and outcomes hence goal-directed decision-making.

According to (Shadab Siddiqui et al., 2025), DRLNDT has better perceptual accuracy with a success rate of up to 99.9 percent in navigation exercises. Moreover, the decision transformer architecture is able to process sequential positional data in parallel, which is much faster and more effective in training in complex environments.

5.4. General Transformer

Table 5.3: Papers that use the General Transformer model

Paper title
Attention-Based Distributional Reinforcement Learning for Safe and Efficient Autonomous Driving (J. Liu et al., 2024)
Balancing Individual and Collective Interests: Integrating Transformer-Based PPO and Social Intelligence for Autonomous Vehicles at Unsignalized Intersections (Doroudian & Taghavifar, 2024)
Empirical Evaluation of Deep Learning based Autonomous Driving System with Real-Time Scene Classification Procedure (Prasath A & Navaz K, 2025)
HGRL: Human-Driving-Data Guided Reinforcement Learning for Autonomous Driving (Zhuang et al., 2024)
Improving policy training for autonomous driving through randomized ensembled double Q-learning with Transformer encoder feature evaluation (Fan et al., 2024)
Multi-Modal Attention Perception for Intelligent Vehicle Navigation Using Deep Reinforcement Learning (Z. Li et al., 2025)
Pre-trained Transformer-Enabled Strategies with Human-Guided Fine-Tuning for End-to-end Navigation of Autonomous Vehicles (Hu et al., 2024)
Safety-constrained Reinforcement Learning with Interaction-aware for Decision-making of Autonomous Driving (Zhang et al., 2025)
SoftCTRL: Soft conservative KL-control of Transformer Reinforcement Learning for Autonomous Driving (Huynh & Nguyen, 2024)
Stability-Aware Reinforcement Learning for Autonomous Driving With Dynamics-Augmented State and Lyapunov Constraints (Luo et al., 2025)
TSAC: Transformer-based Soft Actor-Critic for Behavior Decision-making in Autonomous Driving (Z. Gao et al., 2025)
Autonomous Driving Agents for Safe Road Junctions Through Deep Reinforcement Learning (Yun & Cao, 2024)

The general Transformer architecture is used as the core component of the state encoder or policy network in most studies. These are the models that are built on the typical Encoder-Decoder or encoder-only architecture but tailored to the autonomous driving tasks, e.g., by adding multimodal fusion mechanisms, e.g., CNNs and LSTMs.

These models do not only receive one form of data input, like only visual

data, but instead combine data from a variety of sensors, like vehicle posture and velocity, or lidar data.

(Prasath A & Navaz K, 2025) is a combination of Transformer and CNN models. It accepts three kinds of data as input: Ego State (steering angle, velocity), Navigation Info and Surrounding Info (lidar). Since the three inputs are of varying attributes, the system applies three independent fully connected networks (FCNs) to process the ego state, navigation, and surrounding information respectively, and convert them into feature vectors of the same length (e.g., $L=256$).

5.5. Other Transformers

Table 5.4: Papers that use the Other Transformer model

Paper title	Transformer model
Augmenting Reinforcement Learning With Transformer-Based Scene Representation Learning for Decision-Making of Autonomous Driving (H. Liu et al., 2024)	multi-stage Transformer and sequential latent Transformer
Constraint-Guided Behavior Transformer for Centralized Coordination of Connected and Automated Vehicles at Intersections (Zhao et al., 2024)	Behavior Transformer
Deep Attention Driven Reinforcement Learning (DAD-RL) for Autonomous Decision-Making in Dynamic Environment (Chowdhury et al., 2024)	Spatio-Temporal Attention Encoder
HGRL: Human-Driving-Data Guided Reinforcement Learning for Autonomous Driving (Zhuang et al., 2024)	GPT2
MTD-GPT: A Multi-Task Decision-Making GPT Model for Autonomous Driving at Unsignalized Intersections (Liu et al., 2023)	GPT2

Transformer architectures have also been suggested to be customized or optimized to particular autonomous driving tasks, or other pre-trained language models (including GPT-2) can be used as decision encoders.

Both (Zhuang et al., 2024) and (Liu et al., 2023) used a generative pre-training architecture that is based on GPT-2 to process decision-making in

autonomous driving. MTD-GPT centers on the enhancement of the accuracy of decision-making in intersections without traffic lights using multi-task learning, whereas HGRL additionally uses human expert information and uses the causal modeling power of GPT-2 to maximize the stability and efficiency of reinforcement learning.

To improve the scene representation and accelerate the training process, (H. Liu et al., 2024) applied MST and SLT Transformers. The Behavior Transformer was applied to solve the problem of multi-vehicle conflicts and absence of rule constraints (Zhao et al., 2024). Spatio-Temporal Attention Encoder was used in (Chowdhury et al., 2024) to overcome high computational complexity and spatiotemporal interaction.

Chapter 6: Implementation Plan

6.1. Training Resources

The training process will heavily utilize GPU resources. I will use GPU resources provided by the MMU and the GPU on the local server for training. The MMU's GPU will be primarily used for training and testing, while the local GPU will be used for the CARLA simulator.

6.2. Training Process

To accelerate the agent's training process, I propose a progressive difficulty training method. The agent is first trained in low-difficulty scenarios, where the model learns the basic principles of driving. In medium-difficulty levels, environmental factors such as pedestrians and inclement weather will be added. In high-difficulty levels, I will introduce some unexpected events.

6.3. Model Design

The algorithm will use the relatively safe SAC identified in this literature review. The general Transformer model will serve as the Actor in the SAC algorithm. The model receives state data from a forward-facing camera, LiDAR point cloud, etc., and outputs values such as Steering $[-1.0, 1.0]$, Throttle/Brake $[-1.0, 1.0]$, where less than 0 represents Brake and greater than 0 represents Throttle.

6.4. Vehicle Design

The vehicle design will utilize CARLA's built-in sensors and vehicle components. The sensor suite will employ RGB cameras, LiDAR, GNSS, and IMU sensors. The cameras will be configured at 1280x720 pixels with an FOV of 90 and mounted at the front, left, and right. The LiDAR will be configured with a range of 80 meters and a cloud point rate of 500,000 per second. The vehicle will utilize a Tesla Model 3.

6.5. Scene Selection

FYP2 will target the highest-risk intersection scenario identified in this paper and simulate it within the CARLA environment.

Chapter 7: Conclusion and Outlook

7.1. Summary of Research

This study systematically reviews academic literature from 2019 to 2025, exploring the safety performance of Deep Reinforcement Learning (DRL) algorithms based on the Transformer model in high-risk autonomous driving scenarios.

- Through data extraction and analysis from these papers, it was found that SAC performs relatively safely in most autonomous driving tasks. However, the choice of algorithm largely depends on the characteristics of the task, and there is no absolutely optimal algorithm.
- Through data extraction and logical derivation, a scenario complexity hierarchy model was derived in this study. The results show that highways < (left-turn intersections/intersections) < merging < roundabouts < (right-turn intersections/intersections). This indicates that intersection turning is currently the most challenging scenario for autonomous driving models.
- This literature review found that most studies tend to use a General Transformer as the core state encoder or policy network. This may be because the Transformer model can effectively integrate heterogeneous data such as vehicle attitude, speed, and LiDAR data through its excellent self-attention mechanism.

7.2. Future Outlook

Based on the previous literature review, the next stage of this Final Year Project 2 (FYP2) will shift to actual simulation experiments:

- The simulator will be based on the CARLA simulator, utilizing the GPU resources provided by the MMU for intensive reinforcement learning training. Google Colab is an alternative.

- The algorithm, model, and scenario selection for FYP2 are all based on the research results of the literature review. The algorithm will use SAC, the model will use a general Transformer, and the scenario will be an intersection.
- FYP2 will focus on developing and training an agent integrating the Transformer and SAC algorithms. The goal is to minimize the collision rate of vehicles in the intersection scenario, thereby improving safety.